

---

# 3. Fundamentals of Deep Learning



Perceptron

Feed-forward Neural Networks

Architecture and Learning in Neural Networks


---

# Perceptron are loosely based on neurons



---

# Perceptron




$$h = w_{11}x_1 + w_{12}x_2 + b$$

$$g = \sigma(h)$$

---

# Neural Network




---

# Building Neural Networks

- Frame the problem
    - What are you trying to predict?
    - What are the predictors ?
    - What known data do you have?
  - Architect the network
  - Train the network with known dataset (labeled data)
  - Use the network to predict unseen data
-

# How do you train a Neural Network?

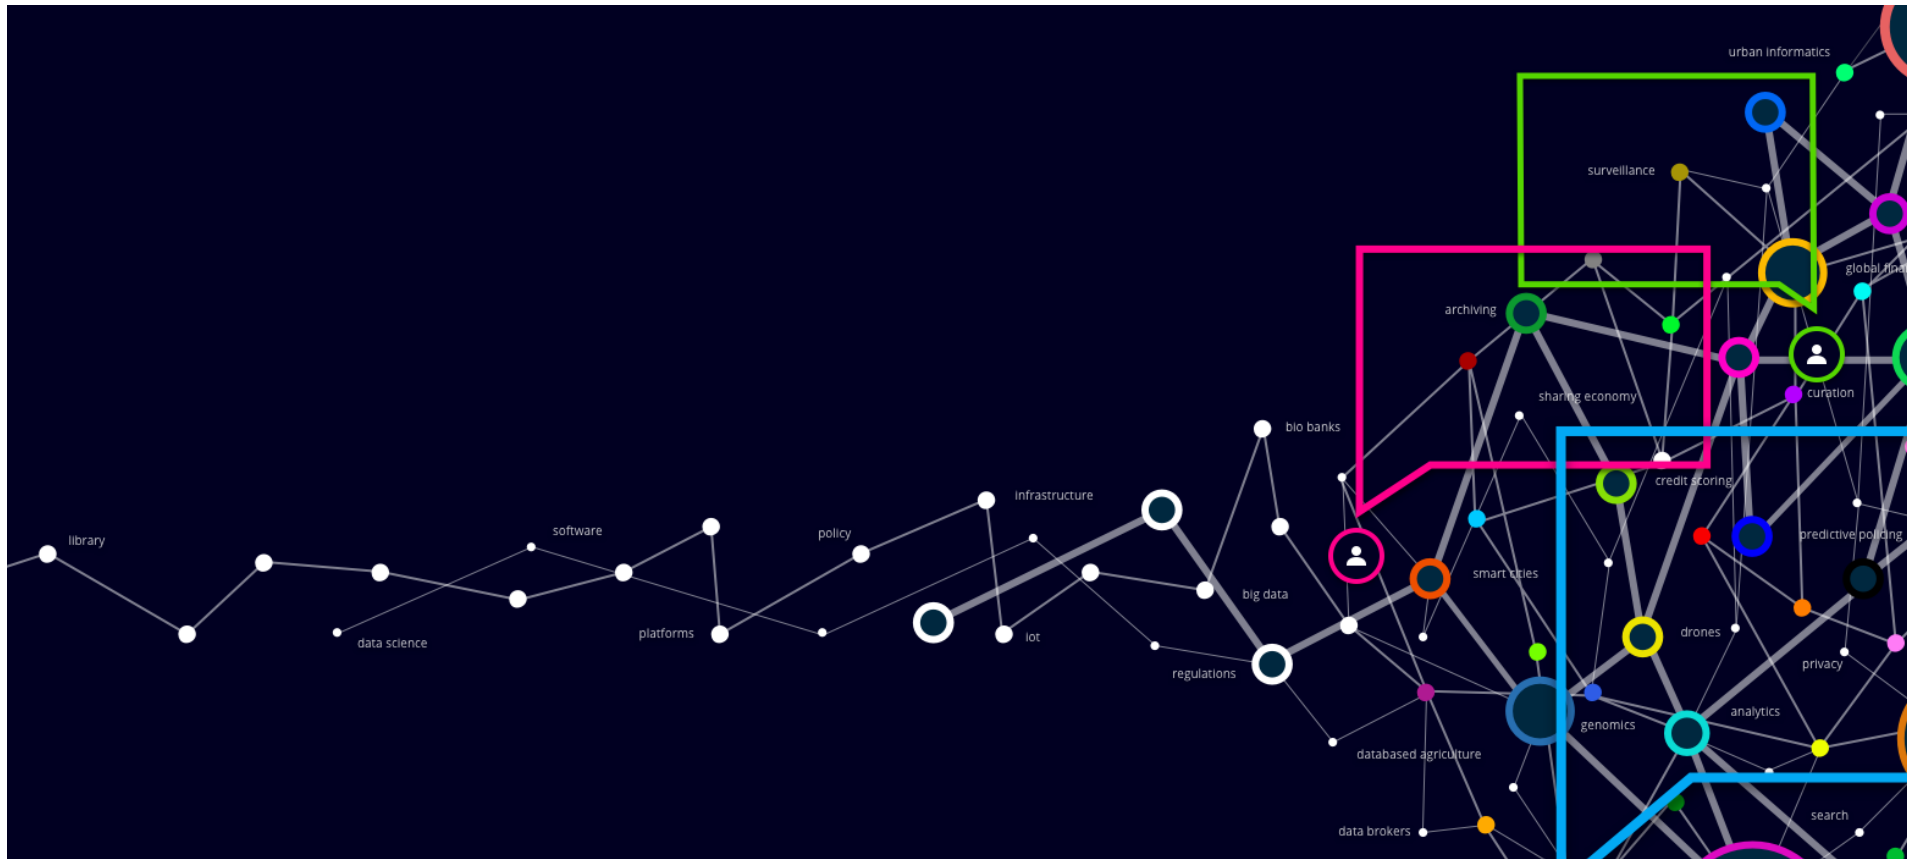


1. **Prepare the data** as input to feed into the neural network

2. **Feedforward** the input through the network, to get the output

3. Compare against the "correct" output,  
**backpropagate** to update weights


4. Repeat



## Preparing Data

---

# Labeled Training Data





---

## Encoding Non-numeric input data

- What if the data has a column "Rainy Day"  
– Y/N.

How would we feed this to the network?

Rainy Day	encoded_rainy_day
Y	1
N	0

---

## Encoding Non-numeric input data

- What if the data has a column "Countries"
  - "Canada", "USA", "Mexico".

How would we feed this to the network?

Country	Country_Canada	Country_USA	Country_Mexico
Canada	1	0	0
USA	0	1	0
Mexico	0	0	1


This encoding is called One-hot encoding

---

---

# Prepare the data

- Rules of thumb
  - Preserve existing relationship between features
  - Do not make nominal data into ordinal
  - Normalize data by scaling using mean/std-dev



## Backpropagation & Training

---

# Backpropagation

- Our objective is to reduce the error.

Assume some random weights

repeat

calculate our model  $y=f(x)$

come up with a way to measure error

adjust weights in order to minimize error

until error (accuracy level) is acceptable


---

# Objective function

- Our approach to measuring error and minimizing it, is called the objective function.
  - Error is measured in many ways.  
This is referred to as **loss function** or **cost function**
  - There some common ones,
    - Mean Square Error (MSE)
    - Mean Absolute Error
    - Logistic Error
    - Cross Entropy
-


---

# Minimizing Error



---

# Classification



$$h = w_{11}x_1 + w_{12}x_2 + b$$


$$g = \sigma(h)$$

---



---

# Classification

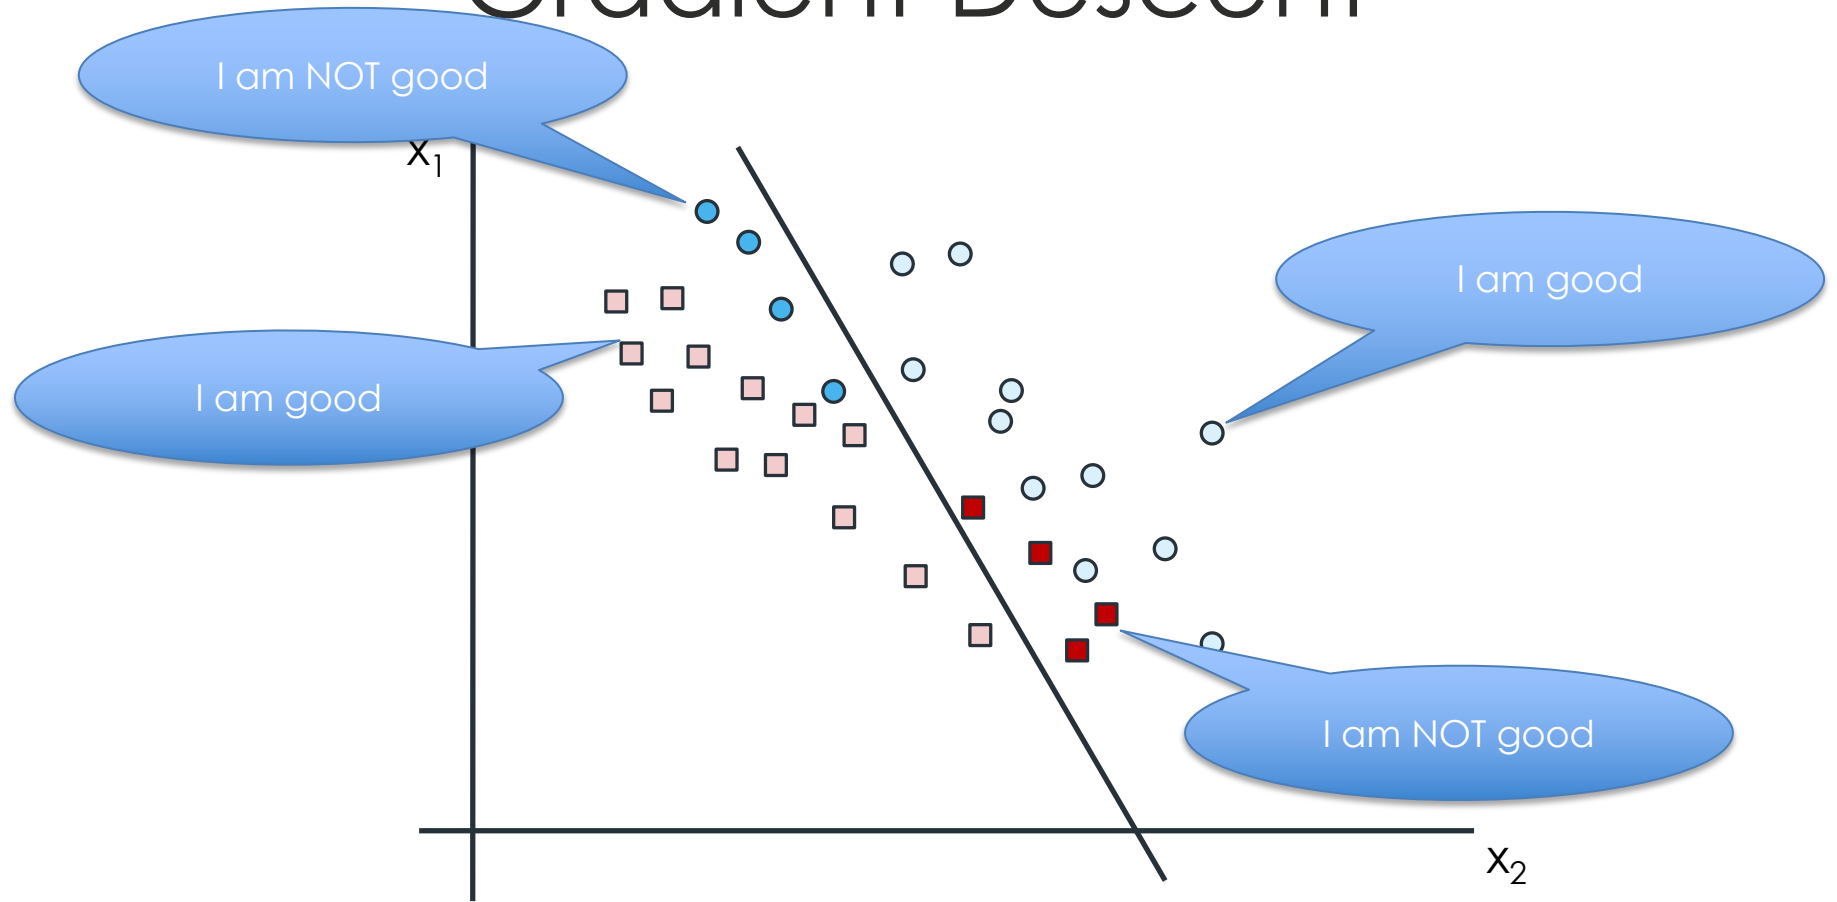


$$h = Wx + b$$

$$g = \sigma(h)$$

---


# Gradient Descent



$$\hat{y} = \sigma(Wx + b)$$

---

# Gradient Descent




$$\hat{y} = \sigma(Wx + b)$$

---

---

# Gradient Descent




$$\hat{y} = \sigma(Wx + b)$$

---

---

# Gradient Descent



$$\hat{y} = \sigma(Wx + b)$$

---

---

This general approach is called Gradient Descent

---

---



# Epochs

- Repeat the training exercise multiple times. Each cycle is called an epoch (pronounced 'epic')
  - Each epoch runs a training cycle that runs in memory. But, the dataset might be too large to fit in memory, so we break each epics into multiple batches. This can be organized as,
    - Batched Gradient Descent
      - Run all the data in each epoch
    - Stochastic Gradient Descent
      - Run random subset of data in each epoch
-

---

# Converting Discrete Errors into Continuous Errors

- Discrete



- Continuous

$P(\text{Yes}) = 85\%$

$P(\text{Red}) = 0.4$   
 $P(\text{Blue}) = 0.5$   
 $P(\text{Green}) = 0.1$




---

# Common Error Functions

- Cross-entropy
  - Good of using with probabilities
  - Works well for classification problems
- Mean-Squared Error
  - Good for continuous predictions


---

# Gradient Descent is a slow process



---


# Adapting Learning Rate



- Make big jumps farther away from minima
- Make smaller (surgical) moves closer to a minima

---

# Adapting with Momentum



- Adding momentum can help the function cross small dips


---

# Common optimizers

- Some common optimizers include,
  - AdamOptimizer
  - RmsPropOptimizer
  - AdaGrad

---

# Avoiding Local Minima




**Random Restarts**

---

---

# Activation - Sigmoid


- Sigmoid – the original Activation function
- Makes the model non-linear
- Vanishing Gradient with Sigmoid – Away from the mean, sigmoid tangents are nearly flat



# Better activation functions




## Hyperbolic Tangents (tanh)

$$\tanh = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



## Rectified Linear Unit (relu)

$$\text{relu}(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{if } x < 0 \end{cases}$$





---

# Regularization

- Penalize Large Weights
- Make the activation function shallow, so GD works better

L1:

$$ErrorFunction = -\frac{1}{m} \sum_{i=1}^n [y_i p_i + (1 - y_i) (1 - p_i)] + \lambda(|w_1| + |w_2| + \dots |w_n|)$$

L2:

$$ErrorFunction = -\frac{1}{m} \sum_{i=1}^n [y_i p_i + (1 - y_i) (1 - p_i)] + \lambda(w_1^2 + w_2^2 + \dots w_n^2)$$

---

---

# Layers

- Dense

Fully Connected Layer

- Dropout

To reduce overfitting. Without dropout, it is likely that the dominant nodes train more than the non-dominant nodes. Dropout gives other nodes a chance. Each node has a probability it will be dropped in a particular epoch during training.

- Flatten

- Convolutional

Apply filters to convert large shallow datasets, into smaller deeper datasets.

- Pooling

Reduce size of dataset, by reducing resolution across local cells

- Locally Connected Layers

- BasicRNNCell

- LSTM ...

---

# Hands-on – Bike Rental Prediction

On workdays, most bikes are rented on warm mornings and evenings

